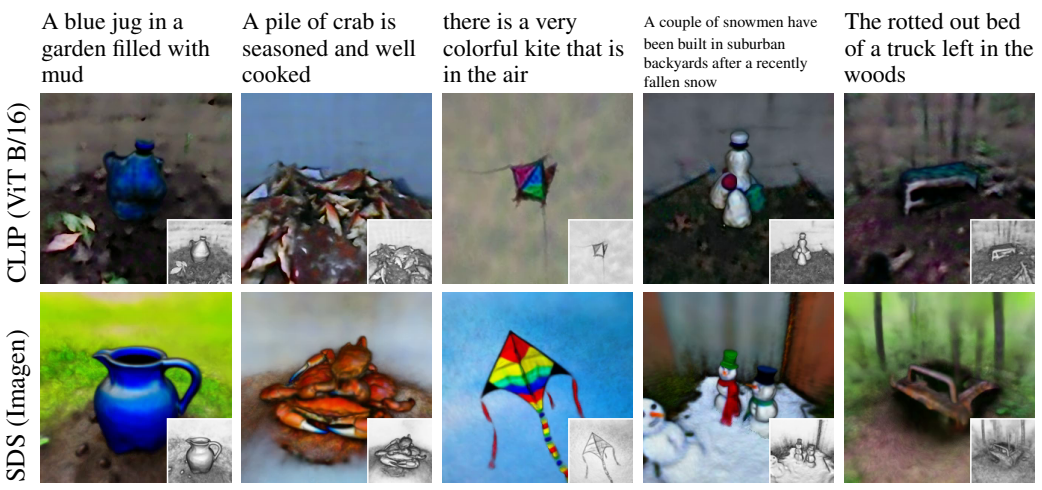

Figure 11: Comparison of the CLIP loss from Dream Fields (top) and our SDS loss. For fair comparison of the loss functions in isolation, we use all of our proposed methods including view-dependent prompts, shading, optimizing untextured renderings, and regularizers, with the same 3D NeRF representation. Qualitatively, 3D scenes generated with SDS (ours) are much more coherent than scenes generated with a CLIP loss.